Bayesian Regression

Hamid Dehghani
School of Computer Science
Birmingham
September 2020

Slides adapted from Iain Styles, School of Computer Science

UNIVERSITY OF BIRMINGHAM

Intended Learning Outcome

- Reason about regression using methods of probability
- Understand how likelihood maximisation and least-squares fitting are related
- Understand the role of prior information in machine learning

UNIVERSITY OF BIRMINGHAM

Using Least Squares Error (LSE)

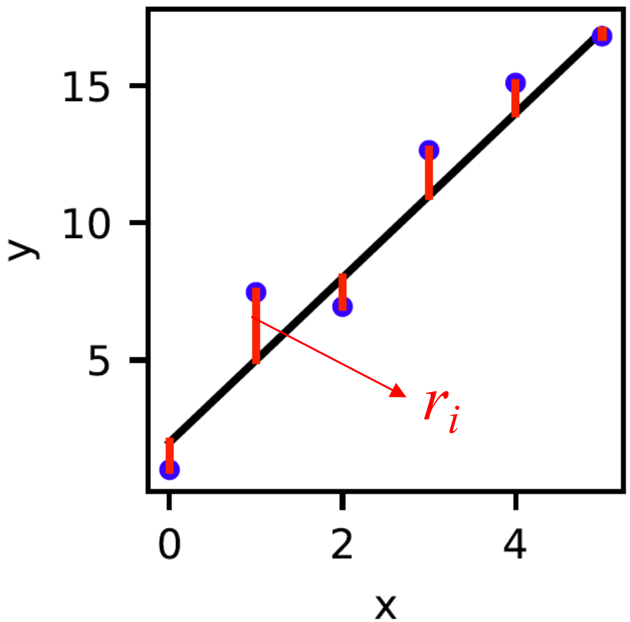

- It is an optimization problem
  - A 'loss/cost' function such that it minimized the difference between measured and modelled data
- Residual $r_i(\mathbf{w})=y_i-f(\mathbf{w},x_i)$

- Why choose this approach?
- Why not some other form of the loss?
- Probabilistic approach will help us understand

UNIVERSITY OF BIRMINGHAM

Modelling the data-generating process

Starting point: model the underlying data-generating process

Assume data points generated by process that has a deterministic component, and some associated sampling uncertainty.

$$y = \mathbf{f}(x, \mathbf{w}) + \epsilon$$

$\epsilon \sim \mathcal{N}(0, \sigma^2)$

$y(x)$ drawn from a normal distribution with mean $f(x, \mathbf{w})$ and variance σ^2

Modelling the data-generating process

We can write the distribution of y as

$$p(y|x, \mathbf{w}, \sigma^2) = \mathcal{N}(y|f(x, \mathbf{w}), \sigma^2)$$

Normal distribution with mean $f(x, \mathbf{w})$, variance σ^2

Note that it is conditional on x, $\mathbf{w}$, and σ

The joint distribution

Dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^{N}$ which we will write as $(\mathbf{x}, \mathbf{y})$.

Assume the y_i are sampled independently normal distributions with the same variance σ^2

Joint PDF is then

$$p(\mathbf{y}|\mathbf{x}, \mathbf{w}, \sigma^2) = \prod_{i=1}^{N} \mathcal{N}(y_i | f(x_i, \mathbf{w}), \sigma^2)$$

The *likelihood* of y

PDF of measurements given parameters

UNIVERSITY OF BIRMINGHAM

Maximum Likelihood

Can now ask "what are the most likely measurements"
Maximise the likelihood
Substitute in the full form of the normal distribution
$\mathcal{N}(x|\mu, \sigma^2) = (2\pi\sigma^2)^{-\frac{1}{2}} \exp(-(x-\mu)^2/(2\sigma^2))$

$$p(\mathbf{y}|\mathbf{x}, \mathbf{w}, \sigma^2) = (2\pi\sigma^2)^{-\frac{N}{2}} \prod_{i=1}^{N} \exp(-(y_i - f(x_i, \mathbf{w}))^2/(2\sigma^2))$$

Take the logarithm (log is monotonic so has same maximum)

$$\ln p(\mathbf{y}|\mathbf{x}, \mathbf{w}, \sigma^2) = \ln(2\pi\sigma^2)^{-\frac{N}{2}} + \ln\left(\prod_{i=1}^{N} \exp(-(y_i - f(x_i, \mathbf{w}))^2/(2\sigma^2))\right)$$

UNIVERSITY OF BIRMINGHAM

Maximum Likelihood

$$\ln p(\mathbf{y}|\mathbf{x}, \mathbf{w}, \sigma^2) = -\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \sum_{i=1}^{N} (y_i - f(x_i, \mathbf{w}))^2$$

- First term (negative) maximised by minimising the number of data points or the variance
- More data and/or more noise means less certainty (accumulation of errors)
- Second term: negative least-squares error
- Maximising the likelihood minimises the least-squares error

Including Priors

Likelihood allows us to apply Bayes rule to include prior knowledge

$$p(a|b) = p(b|a)p(a)/p(b)$$

$p(a|b)$ is the posterior distribution of a given b, $p(b|a)$ is the likelihood of b given a and $p(a)$ is the prior distribution of a.

Can now ask: given a set of measurements, how are the weights distributed?

$$p(\mathbf{w}|\mathbf{x},\mathbf{y},\sigma^2) = \frac{p(\mathbf{y}|\mathbf{x},\mathbf{w},\sigma^2) \times p(\mathbf{w})}{P(\mathbf{y})}$$

Ignore $P(\mathbf{y})$ for simplicity

$$p(\mathbf{w}|\mathbf{x},\mathbf{y},\sigma^2) \propto p(\mathbf{y}|\mathbf{x},\mathbf{w},\sigma^2) \times p(\mathbf{w})$$

Including Priors

Consider $p(\mathbf{w}) = c$, a constant.

All parameter values equally likely

$$\begin{aligned} p(\mathbf{w}|\mathbf{x}, \boldsymbol{y}, \sigma^2) &\propto p(\mathbf{y}|\mathbf{x}, \mathbf{w}, \sigma^2) \times c \\ &\propto p(\mathbf{y}|\mathbf{x}, \mathbf{w}, \sigma^2) \end{aligned}$$

The same max likelihood problem as before

The least squares error assigns model weights that are uniformly distributed

Is this desireable?

UNIVERSITY OF BIRMINGHAM

Including Priors

- Uniform distribution of weights seems reasonable
- But allows very large high-frequency terms to match model noise

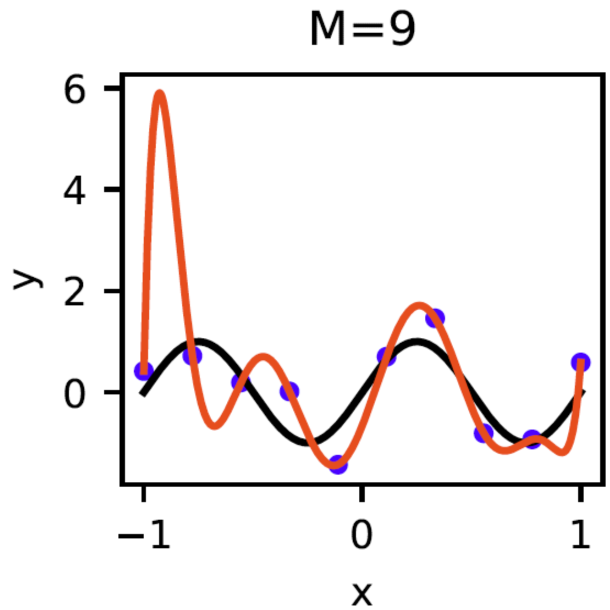

M	w_0	w_1	w_2	w_3	w_4	w_5	w_6	w_7	w_8	w_9
9	-0.66	10.98	25.62	-117.80	-143.29	405.10	246.74	-561.32	-127.91	263.129

UNIVERSITY OF BIRMINGHAM

Gaussian Prior

How to make large weights unlikely?

Gaussian prior: most weights near zero

$$
\begin{aligned}
p(\mathbf{w}|\lambda) &\propto \prod_{i=1}^{M} \exp(-\lambda w_i^2) \\
&\propto \exp(-\lambda \sum_i w_i^2) \\
&\propto \exp(-\lambda \mathbf{w}^{\mathrm{T}} \mathbf{w})
\end{aligned}
$$

Conditioned on parameters $\lambda = 1/2\sigma^2$ (ie large lambda *mapsto* narrow distribution)

UNIVERSITY OF BIRMINGHAM

Summary

- Probabilistic formulation of regression
- Maximising likelihood minimises least squares error
- Prior distributions of parameters
- Reading: Bishop, section 1.2.5