




k-Nearest Neighbours


Leandro L. Minku

k-Nearest Neighbours: Basic Idea




$y \in \{\text{blue, orange}\}$

k-Nearest Neighbours: Basic Idea




Usually, for classification problems: predict the majority among the values of the dependent variable of the k nearest neighbours (majority vote).

k-Nearest Neighbours: Basic Idea




Usually, for classification problems: predict the majority among the values of the dependent variable of the k nearest neighbours (majority vote).

k-Nearest Neighbours: Basic Idea




What is the predicted output for this test instance?

k-Nearest Neighbours: Basic Idea




What is the predicted output for this test instance?

k-Nearest Neighbours: Basic Idea



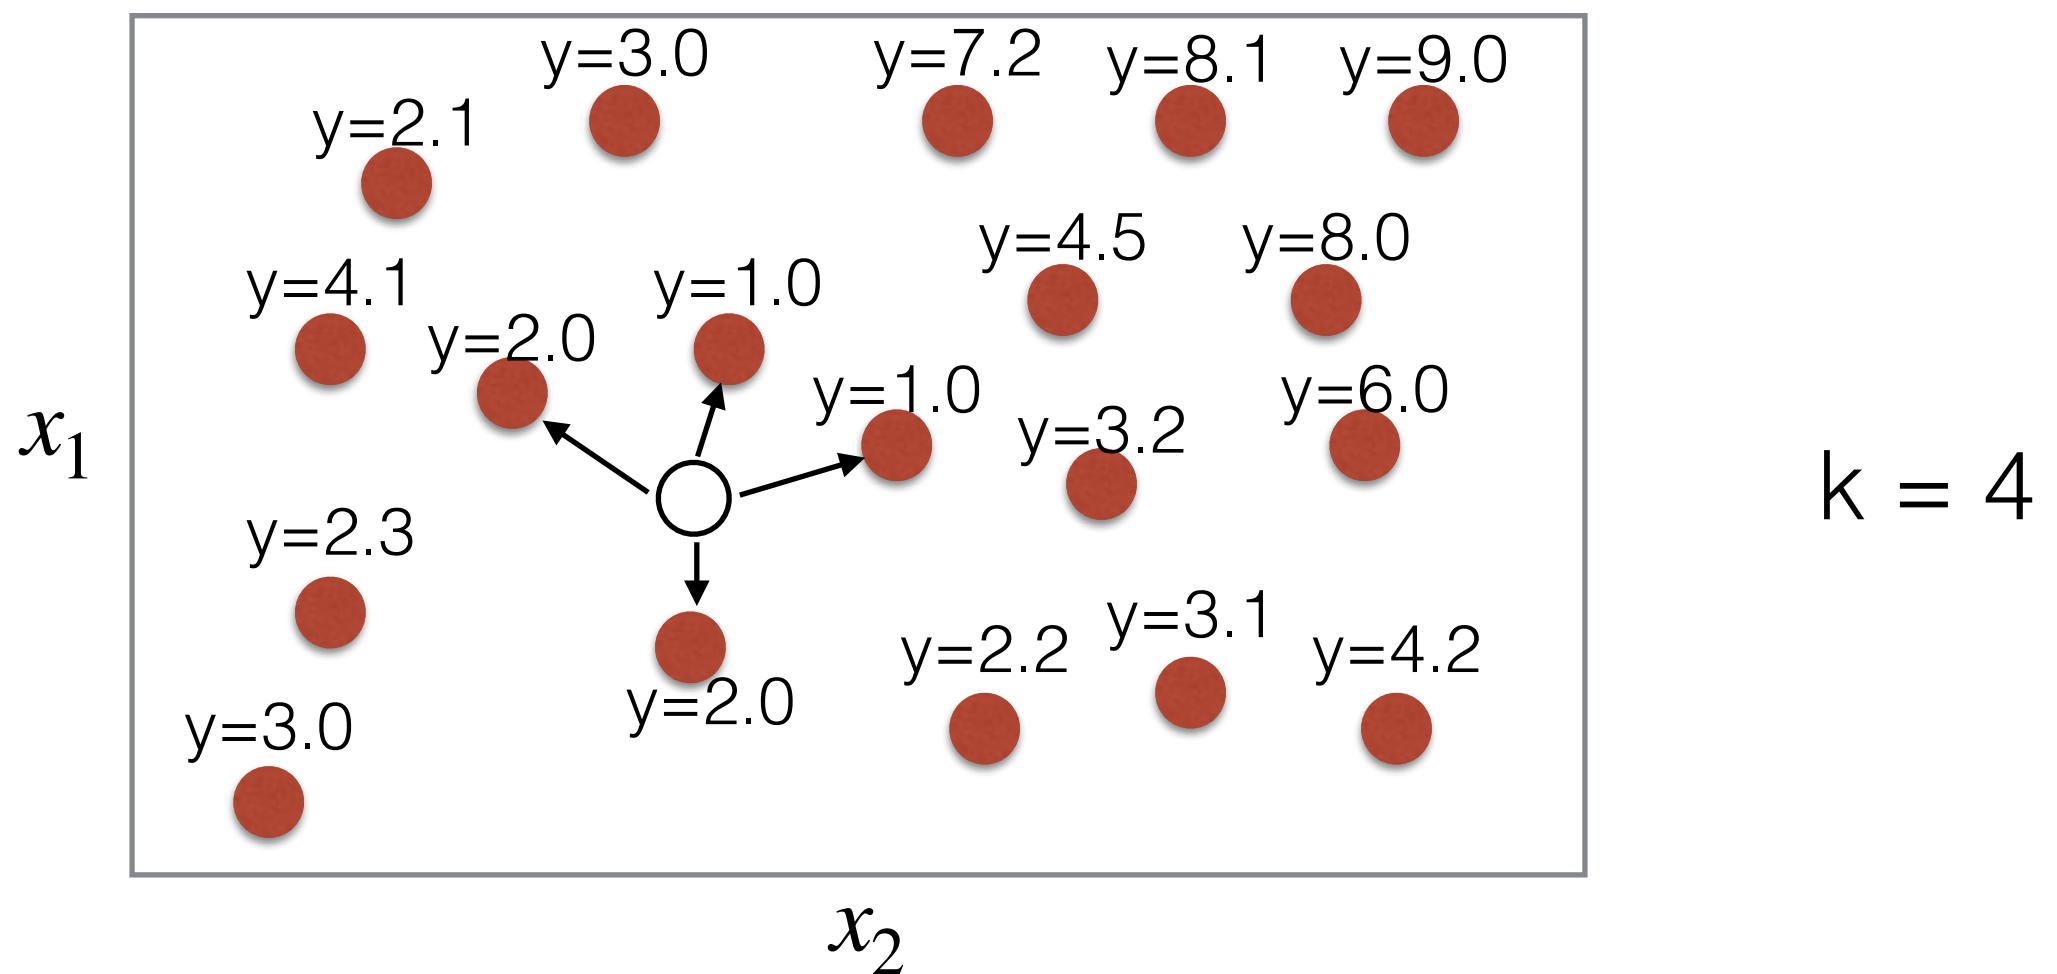
What is the predicted output for this test instance?

k-Nearest Neighbours: Basic Idea




What is the predicted output for this test instance?

k-Nearest Neighbours: Basic Idea



Usually, for regression problems: predict the average among the values of the dependent variable of the k nearest neighbours.

k-Nearest Neighbours: Basic Idea



Given an instance to be predicted, we need to find its k nearest neighbours, based on some distance metric **on the input space**.

Distance Metric

- Usually, this is the Euclidean Distance.
- For d dimensions in the input space:

$$\text{distance}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \sqrt{(x_1^{(i)} - x_1^{(j)})^2 + (x_2^{(i)} - x_2^{(j)})^2 + \dots + (x_d^{(i)} - x_d^{(j)})^2}$$

$$\text{distance}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \sqrt{\sum_{p=1}^d (x_p^{(i)} - x_p^{(j)})^2} = \sqrt{(\mathbf{x}^{(i)} - \mathbf{x}^{(j)})^T (\mathbf{x}^{(i)} - \mathbf{x}^{(j)})}$$

Normalisation of Numeric Independent Variable

- **Problem:** different numeric independent variable may have different scales.
 - Scale of numeric independent variable will influence the Euclidean Distance.
 - If x_1 is in $[0, 10]$ and x_2 is in $[100, 10000]$, x_2 will influence the distance more.

- **Popular solution:**

- Normalise numeric independent variables of all data so that they will be between 0 and 1. E.g.: normalising independent variable p of example i :

$$\text{normalise}(x_p^{(i)}) = \frac{x_p^{(i)} - \min_p}{\max_p - \min_p}$$

- **How to know the minimum and maximum values?**

- If the real minimum and maximum are unknown, for each input attribute, use the minimum and maximum values present in the training set.

Ordinal or Categorical Independent variable

- Independent variable can be numerical, ordinal or categorical.
 - **Numeric**: e.g., age, salary.
 - **Ordinal**: e.g., expertise in {low, medium, high}.
 - **Categorical**: e.g., car in {fiat, volkswagen, toyota}.

- Euclidean distance is defined for numerical data!

$$\text{distance}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \sqrt{\sum_{p=1}^d (x_p^{(i)} - x_p^{(j)})^2}$$

- For ordinal independent variable, we can convert them to numeric.
 - E.g.: low = 0, medium = 0.5, high = 1.
- For categorical independent variable, we could use the following idea:

$$\begin{aligned} &\text{if } (x_p^{(i)} = x_p^{(j)}), \quad (x_p^{(i)} - x_p^{(j)}) = 0 \\ &\text{if } (x_p^{(i)} \neq x_p^{(j)}), \quad (x_p^{(i)} - x_p^{(j)}) = 1 \end{aligned}$$

Procedure for Predicting a New (Test) Example

- Given the independent variables of a new example $(\mathbf{x}^{(i)}, ?)$, the number of neighbours k , and **min** and **max** observed so far.
- Update **min** and **max** based on $\mathbf{x}^{(i)}$
- For each training example $(\mathbf{x}^{(j)}, y^{(j)})$
 - $\text{dist} = \text{distance}(\text{normaliseEachVar}(\mathbf{x}^{(i)}, \mathbf{min}, \mathbf{max}), \text{normaliseEachVar}(\mathbf{x}^{(j)}, \mathbf{min}, \mathbf{max}))$
 - Add $(\text{dist}, y^{(j)})$ to a data structure T sorted based on ascending order of distance.
- Return the majority vote (or average) of $y^{(j)}$ for the first k entries of T .

k-NN Approach

- k-NN Learning Algorithm:
 - No real training; simply store all training data received so far, together with the maximum and minimum values of the numerical independent variables.
- k-NN “Model”:
 - All training data received so far, together with the maximum and minimum values of the numerical independent variables.
- k-NN prediction for an instance ($\mathbf{x}^{(i)}, ?$):
 - Find the K nearest neighbours, i.e., the K training examples that are the closest to $\mathbf{x}^{(i)}$.
 - For classification problems: majority vote.
 - For regression problems: average.

Advantages and Disadvantages

- Advantages:
 - Training is simple and quick: store the training data.
- Disadvantage:
 - Memory requirements are high: stores all data, which can be troublesome when training set is very large.
 - Making predictions is slow: we have to search for the nearest neighbours among all the training data, which can be troublesome when training set is very large.

[Original] k-NN is not adequate when we have very large training sets.

Advantages and Disadvantages

- Advantages:
 - Training is simple and quick: store the training data.
- Disadvantage:
 - Memory requirements are high: stores all data, which can be troublesome when training set is very large.
 - Making predictions is slow: we have to search for the nearest neighbours among all the training data, which can be troublesome when training set is very large.

k-NN can be good for applications where there is little data.

Advantages and Disadvantages

- Advantages:
 - Training is simple and quick: store the training data.
- Disadvantage:
 - Memory requirements are high: stores all data, which can be troublesome when training set is very large.
 - Making predictions is slow: we have to search for the nearest neighbours among all the training data, which can be troublesome when training set is very large.

Intuitive: k-NN helps people to find the examples that are most similar to the new example.

Quiz

- Consider that you have an example with $\mathbf{x}^{(1)} = (5, 10)^T$. Consider also that $\min_1 = 0$, $\max_1 = 10$, $\min_2 = 0$, $\max_2 = 20$. What is the normalised value of $\mathbf{x}^{(1)}$?

- Consider a given training set containing the following normalised examples:

- $\mathbf{x}^{(1)} = (0.1, 0.1)^T$, $y^{(1)} = \text{red}$
- $\mathbf{x}^{(2)} = (0.1, 0.2)^T$, $y^{(2)} = \text{blue}$
- $\mathbf{x}^{(3)} = (0.2, 0.2)^T$, $y^{(3)} = \text{green}$

What class would be predicted for a normalised test example $\mathbf{x}^{(4)} = (0.3, 0.3)^T$ when $k=1$?

Further Reading

- Essential:

- Iain Style's notes on "Classification and k-Nearest Neighbours".

- Recommended:

- Section 2.5.2 of Bishop, Pattern Recognition and Machine Learning, contains a brief treatment of nearest-neighbour methods.
- A Detailed Introduction to K-Nearest Neighbor (KNN) Algorithm by Saravanan Thirumuruganathan. Access at: <https://saravananthirumuruganathan.wordpress.com/2010/05/17/a-detailed-introduction-to-k-nearest-neighbor-knn-algorithm>

- Suggested:

- Russell and Norvig's "Artificial Intelligence: A Modern Approach"
 - Section 18 (Learning from Examples) up to the end of section 18.2 (Supervised Learning).
 - Section 18.8 (Non-Parametric Models) up to the end of section 18.8.1 (Nearest Neighbour Models).